

# Artificial Intelligence – Semester Project Guidelines

## 1. Objective

The purpose of this semester project is to provide students with an opportunity to apply **Artificial Intelligence (AI) concepts, algorithms, and design methodologies** to solve real-world or cybersecurity-related problems.

Students will design, implement, and evaluate a complete **AI-driven system**, integrating:

- **Search algorithms,**
- **Constraint satisfaction and optimization,**
- **Supervised and unsupervised learning, and**
- **Reinforcement learning** concepts.

The project will mirror **industry-level workflows**, emphasizing data pipelines, reproducibility, explainability, and AI ethics.

## 2. Learning Outcomes

By the end of this project, students will be able to:

- Formulate AI problems with measurable goals and evaluation metrics.
- Implement search and learning algorithms efficiently.
- Handle real-world datasets (preprocessing, feature engineering, scaling).
- Develop and evaluate ML and RL models using modern libraries.
- Apply explainable AI techniques (SHAP, LIME).
- Document and present AI-driven solutions using professional reporting standards.

## 3. Project Lifecycle and Weekly Integration

Weeks	Course Topics	Project Phase / Expected Deliverables
Week 7	<i>Introduction to AI &amp; Intelligent Agents</i>	Define problem domain, objectives, and agent type (reflex, model-based, goal-based, or learning agent). Draft <b>Project Proposal</b> .
Week 8	<i>Search Strategies (Uninformed &amp; Informed)</i>	Implement initial search/optimization-based solution (e.g., BFS, DFS, A*, Best-First). Prepare <b>Agent Design Report</b> .
Week 9	<i>Constraint Satisfaction Problems (CSPs)</i>	Model part of your system as a CSP (variables, domains, constraints). Implement backtracking or arc consistency.
Week 10	<i>Machine Learning Foundations &amp; Supervised Learning</i>	Acquire/prepare dataset. Apply data preprocessing, feature selection, and train at least two supervised models (e.g., Decision Tree, Logistic Regression, Random Forest).

<b>Week 11</b>	<i>Advanced ML (SVMs, Neural Networks, CNNs)</i>	Integrate advanced ML or deep learning model. Perform hyperparameter tuning and visualize learning results.
<b>Week 12</b>	<i>Unsupervised Learning &amp; Reinforcement Learning</i>	Extend system using clustering or reinforcement learning (Q-learning, DQN). Analyze model behavior on unseen or adversarial data.
<b>Week 13</b>	<i>Recap &amp; Final Presentation</i>	Integrate components, finalize report, and deliver <b>presentation + demo</b> .

## 4. Project Categories (Select One Domain)

### A. Cybersecurity Intelligence

- **Examples:** Intrusion Detection, Phishing Detection, Network Traffic Analysis
- **Techniques:** Random Forest, SVM, ANN, CNN, Autoencoders
- **Datasets:** UNSW-NB15, CICIDS2017, NSL-KDD
- **Metrics:** Precision, Recall, F1-score, ROC-AUC

### B. Optimization & Search Systems

- **Examples:** Routing optimization, resource allocation, scheduling
- **Techniques:** A\*, Hill-Climbing, Genetic Algorithms, CSPs
- **Evaluation:** Cost minimization, time efficiency, convergence

### C. Intelligent Agent Simulations

- **Examples:** Multi-agent security systems, autonomous defenders, logistics bots
- **Techniques:** Model-based, goal-based, and learning agents
- **Evaluation:** Decision accuracy, adaptability, and computational efficiency

### D. Predictive Analytics

- **Examples:** Threat prediction, log-based anomaly detection, behavior forecasting
- **Techniques:** Regression, Ensemble Learning, Time Series (LSTM, Prophet)
- **Evaluation:** RMSE, MAE, R<sup>2</sup>, trend accuracy

### E. Reinforcement Learning Applications

- **Examples:** Adaptive intrusion response, defense automation, game-based control
- **Techniques:** Q-learning, SARSA, DQN, PPO
- **Evaluation:** Reward convergence, policy stability, exploration balance

### F. NLP & Threat Intelligence

- **Examples:** Text classification, malicious domain detection, sentiment analysis
- **Techniques:** TF-IDF, Word2Vec, BERT, Logistic Regression

- **Datasets:** Phishing Email Corpus, URLNet
- **Evaluation:** F1, ROC-AUC, confusion matrix

## 5. Deliverables and Timeline

<b>Deliverable</b>	<b>Description</b>	<b>Due Week</b>
<b>1. Proposal Report (4–5 pages)</b>	Problem statement, motivation, objectives, dataset source, AI approach, and evaluation plan.	Week 7
<b>2. Progress Report I</b>	Agent/search algorithm implementation with working example.	Week 9
<b>3. Progress Report II</b>	Data preprocessing pipeline + baseline ML models with performance metrics.	Week 8
<b>4. Progress Report III</b>	Integration of advanced ML/DL or RL model with interpretability and optimization.	Week 10
<b>5. Final Report + Demo + Presentation</b>	+ Full working system, final evaluation, report, slides, and GitHub repository.	Week 13

## 6. Technical Standards & Tools

<b>Component</b>	<b>Requirements / Recommendations</b>
<b>Language</b>	Python (mandatory for ML/RL tasks)
<b>Libraries</b>	NumPy, Pandas, scikit-learn, TensorFlow, PyTorch, OpenAI Gym, XGBoost
<b>Data Tools</b>	Kaggle, UCI, UNSW-NB15, CICIDS2017
<b>Visualization</b>	Matplotlib, Seaborn, TensorBoard, Plotly
<b>Explainability Tools</b>	SHAP, LIME
<b>Version Control</b>	GitHub/GitLab (with commit logs)
<b>Experiment Tracking (optional)</b>	Weights & Biases (W&B), MLflow
<b>Deployment (optional)</b>	Streamlit, Flask, or Docker

## 7. Evaluation Rubrics (Detailed and Technical)

<b>Evaluation Component</b>	<b>Description</b>	<b>Weight</b>
<b>1. Problem Definition &amp; Relevance</b>	& Problem clarity, industry linkage, measurable objectives	10%
<b>2. System Architecture Design</b>	& Modular pipeline (data → model → evaluation), scalability, architecture diagram	15%

<b>3. Implementation</b>	<b>Algorithmic</b> Correct and efficient implementation of AI algorithms (search, CSP, ML, RL)	20%
<b>4. Data Handling &amp; Feature Engineering</b>	Data quality, preprocessing, normalization, feature extraction, ethics	15%
<b>5. Model Development &amp; Evaluation</b>	& Model selection, parameter tuning, comparative analysis, validation (e.g., K-fold CV)	15%
<b>6. Interpretability Explainability</b>	& Use of SHAP/LIME, visualization of model insights, explainable decision-making	10%
<b>7. Documentation Presentation</b>	& Report quality, formatting, technical writing, demo clarity	10%
<b>8. Scalability &amp; Innovation</b>	Creativity, reusability, performance optimization, or deployment	5%
<b>Total</b>		<b>100%</b>

## 8. Detailed Evaluation Matrix

Criteria	Excellent (A)	Good (B)	Fair (C)	Poor (D/F)
<b>Problem Formulation</b>	Domain-specific, measurable, innovative	Clear but lacks technical novelty	Vague objectives	Unclear or trivial
<b>Algorithmic Design</b>	Correct, optimized, modular	Correct but limited optimization	Partial or inefficient	Incorrect implementation
<b>Data Engineering</b>	Robust pipeline, visual EDA, engineered features	Good preprocessing	Minimal effort	Raw/unprocessed
<b>Model Evaluation</b>	Multiple metrics, proper validation	Partial evaluation	Basic accuracy only	No evaluation
<b>Explainability</b>	SHAP/LIME visualization used	or Partial interpretation	Minimal effort	None
<b>Report &amp; Presentation</b>	Industry-grade structure, clarity	visually organized	Limited structure	Disorganized

## 9. Reporting Format

Each report (especially final) must follow a structured **academic–industrial hybrid** format:

1. **Abstract**
2. **Introduction & Motivation**
3. **Problem Definition & Objectives**
4. **Literature Review**

5. **System Architecture** (with diagram)
6. **Data Description & Preprocessing**
7. **Algorithmic Implementation** (Search / ML / RL / CSP)
8. **Model Evaluation & Comparison**
9. **Explainability & Visualization**
10. **Results & Discussion**
11. **Ethical AI & Limitations**
12. **Conclusion & Future Work**
13. **References (IEEE format)**

## 10. Technical Best Practices

Aspect	Requirement
<b>Code Quality</b>	Follow PEP8; use docstrings and modular functions.
<b>Reproducibility</b>	Include <code>requirements.txt</code> or <code>environment.yml</code> .
<b>Version Control</b>	Use Git; frequent commits with messages.
<b>Experiment Tracking</b>	Log results, models, and parameters.
<b>Visualization</b>	Include model performance plots (loss, accuracy, ROC).
<b>AI Ethics</b>	Avoid biased datasets; justify all data sources.

## 11. Submission Checklist

1. Proposal (Week 3)
2. Progress Reports I–III (Weeks 5, 8, 12)
3. Final Report (Week 13)
4. Presentation Slides + Demo
5. GitHub Repository (Code + README + `requirements.txt`)
6. (Optional) Streamlit/Flask Deployment or MLflow Logging

## 12. Bonus Marks (Up to +10%)

Criterion	Bonus
Integration of Explainability (SHAP/LIME)	+2%
ML Experiment Tracking (MLflow/W&B)	+3%
Deployed Prototype (Flask/Streamlit/Docker)	+5%

## 13. Example Real-World Project Ideas

Domain	Project Title	AI Techniques
Cybersecurity	Intrusion Detection using Random Ensemble Learning, Explainable Forest & SHAP	AI
Cyber Defense	Q-Learning based Intrusion Response Agent	Reinforcement Learning
Optimization	Hybrid A* + CSP Resource Allocation	Heuristic Search, Constraint Propagation
NLP Threat Intel	Phishing Email Detection using BERT Transformer, Text Classification	
Predictive Analysis	Time-Series Attack Forecasting	LSTM, ARIMA
Autonomous Systems	Multi-Agent Security Simulation	Goal-Based & Learning Agents

## 14. Final Presentation

- **Duration:** 12–15 minutes per team
- **Format:**
  1. Problem introduction & motivation
  2. System architecture & algorithms
  3. Data pipeline & results
  4. Visualization and interpretability results
  5. Demo (if applicable)
  6. Q&A

Evaluation will focus on **clarity, technical mastery, and result justification**.

## 15. Group Policy

- **Group Size:** 2–3 members.
- Each member must contribute technically (e.g., ML module, visualization, RL model, or documentation).
- **Individual performance** verified via contribution logs, commits, and oral defense.